

lightgbm

March 6, 2021

0.1 Problem Statement

0.1.1 Topic: Health Insurance Lead Prediction

Your Client FinMan is a financial services company that provides various financial services like loan, investment funds, insurance etc. to its customers. FinMan wishes to cross-sell health insurance to the existing customers who may or may not hold insurance policies with the company. The company recommend health insurance to it's customers based on their profile once these customers land on the website. Customers might browse the recommended health insurance policy and consequently fill up a form to apply. **When these customers fill-up the form, their Response towards the policy is considered positive and they are classified as a lead.**

Once these leads are acquired, the sales advisors approach them to convert and thus the company can sell proposed health insurance to these leads in a more efficient manner. Now the company needs your help in building a model to predict whether the person will be interested in their proposed Health plan/policy given the information about:

- Demographics (city, age, region etc.)
- Information regarding holding policies of the customer
- Recommended Policy Information

```
[1]: #loading packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: #Reading data
train = pd.read_csv(r'C:\Users\Ranjeet shrivastav\Videos\ML_algorithms\Health_Insurance Lead Prediction\train.csv')
test = pd.read_csv(r'C:\Users\Ranjeet shrivastav\Videos\ML_algorithms\Health_Insurance Lead Prediction\test.csv')
```

```
[3]: train.shape, test.shape
```

```
[3]: ((50882, 14), (21805, 13))
```

```
[4]: train.head(2)
```

```
[4]:   ID City_Code  Region_Code Accomodation_Type Reco_Insurance_Type Upper_Age \
0    1         C3         3213          Rented          Individual        36
1    2         C5         1117           Owned           Joint         75

   Lower_Age Is_Spouse Health Indicator Holding_Policy_Duration \
0         36        No          X1             14+
1         22        No          X2             NaN

   Holding_Policy_Type  Reco_Policy_Cat  Reco_Policy_Premium  Response
0                3.0             22          11628.0           0
1                NaN             22          30510.0           0
```

Its generally a good idea to combine both train and test data sets into one, perform feature engineering and then divide them later again. This saves the trouble of performing the same steps twice on test and train. Lets combine them into a dataframe 'data' with a 'source' column specifying where each observation belongs.

```
[5]: train['source'] = 'train'
test['source'] = 'test'
data = pd.concat([train,test],ignore_index = True)
```

```
[6]: data.head()
```

```
[6]:   Accomodation_Type City_Code Health Indicator Holding_Policy_Duration \
0          Rented        C3          X1             14+
1          Owned        C5          X2             NaN
2          Owned        C5          NaN             1.0
3          Owned       C24          X1             14+
4          Rented        C8          X2             3.0

   Holding_Policy_Type  ID Is_Spouse  Lower_Age Reco_Insurance_Type \
0                3.0    1        No         36          Individual
1                NaN    2        No         22             Joint
2                1.0    3        No         32          Individual
3                3.0    4        No         48             Joint
4                1.0    5        No         44          Individual

   Reco_Policy_Cat  Reco_Policy_Premium  Region_Code  Response  Upper_Age \
0                22          11628.0         3213         0.0         36
1                22          30510.0         1117         0.0         75
2                19           7450.0         3732         1.0         32
3                19          17780.0         4378         0.0         52
4                16          10404.0         2190         0.0         44

   source
0  train
```

```
1 train
2 train
3 train
4 train
```

```
[7]: data.shape
```

```
[7]: (72687, 15)
```

```
[8]: data.dtypes
```

```
[8]: Accomodation_Type      object
     City_Code             object
     Health Indicator       object
     Holding_Policy_Duration object
     Holding_Policy_Type    float64
     ID                    int64
     Is_Spouse              object
     Lower_Age              int64
     Reco_Insurance_Type    object
     Reco_Policy_Cat        int64
     Reco_Policy_Premium    float64
     Region_Code            int64
     Response               float64
     Upper_Age              int64
     source                 object
     dtype: object
```

```
[9]: data.isnull().sum()
```

```
[9]: Accomodation_Type      0
     City_Code             0
     Health Indicator      16718
     Holding_Policy_Duration 28854
     Holding_Policy_Type    28854
     ID                   0
     Is_Spouse             0
     Lower_Age             0
     Reco_Insurance_Type    0
     Reco_Policy_Cat        0
     Reco_Policy_Premium    0
     Region_Code            0
     Response               21805
     Upper_Age              0
     source                 0
     dtype: int64
```

Note that the 'Response' is the target variable and missing values are ones in the test set. So

we need not worry about it. But we'll impute the missing values in 'Health Indicator', 'Holding_Policy_Duration' and 'Holding_Policy_Type' in the data cleaning section.

```
[10]: data.nunique()
```

```
[10]: Accomodation_Type      2
      City_Code             36
      Health_Indicator       9
      Holding_Policy_Duration 15
      Holding_Policy_Type    4
      ID                   72687
      Is_Spouse             2
      Lower_Age             60
      Reco_Insurance_Type    2
      Reco_Policy_Cat        22
      Reco_Policy_Premium    7685
      Region_Code           5538
      Response              2
      Upper_Age             58
      source                2
      dtype: int64
```

```
[11]: data['Health Indicator'].value_counts()
```

```
[11]: X1      18624
      X2      14848
      X3       9608
      X4       8185
      X5       2408
      X6       1794
      X7        292
      X8        119
      X9         91
      Name: Health Indicator, dtype: int64
```

```
[12]: data['Health Indicator'].replace(to_replace='X1', value='0', regex=True,
      ↪ inplace=True)
      data['Health Indicator'].replace(to_replace='X2', value='1', regex=True,
      ↪ inplace=True)
      data['Health Indicator'].replace(to_replace='X3', value='2', regex=True,
      ↪ inplace=True)
      data['Health Indicator'].replace(to_replace='X4', value='3', regex=True,
      ↪ inplace=True)
      data['Health Indicator'].replace(to_replace='X5', value='4', regex=True,
      ↪ inplace=True)
      data['Health Indicator'].replace(to_replace='X6', value='5', regex=True,
      ↪ inplace=True)
```

```
data['Health Indicator'].replace(to_replace='X7', value='6', regex=True,
    ↪inplace=True)
data['Health Indicator'].replace(to_replace='X8', value='7', regex=True,
    ↪inplace=True)
data['Health Indicator'].replace(to_replace='X9', value='8', regex=True,
    ↪inplace=True)
```

```
[13]: data['Holding_Policy_Duration'].value_counts()
```

```
[13]: 1.0      6390
      14+     6227
      2.0     6032
      3.0     5192
      4.0     3976
      5.0     3354
      6.0     2797
      7.0     2309
      8.0     1885
      9.0     1607
     10.0     1146
     11.0      800
     13.0      732
     12.0      709
     14.0      677
      Name: Holding_Policy_Duration, dtype: int64
```

```
[14]: data['Holding_Policy_Duration'] = data['Holding_Policy_Duration'].
    ↪replace('14+', '14.0')
```

```
[15]: data['Is_Spouse'].value_counts()
```

```
[15]: No      60687
      Yes     12000
      Name: Is_Spouse, dtype: int64
```

```
[16]: data['Is_Spouse'] = data['Is_Spouse'].replace('Yes',1)
      data['Is_Spouse'] = data['Is_Spouse'].replace('No',0)
```

```
[17]: data.head()
```

```
[17]: Accomodation_Type  City_Code  Health Indicator  Holding_Policy_Duration  \
0          Rented      C3              0              14.0
1          Owned      C5              1              NaN
2          Owned      C5             NaN              1.0
3          Owned     C24              0              14.0
4          Rented      C8              1              3.0
```

	Holding_Policy_Type	ID	Is_Spouse	Lower_Age	Reco_Insurance_Type	\
0	3.0	1	0	36	Individual	
1	NaN	2	0	22	Joint	
2	1.0	3	0	32	Individual	
3	3.0	4	0	48	Joint	
4	1.0	5	0	44	Individual	

	Reco_Policy_Cat	Reco_Policy_Premium	Region_Code	Response	Upper_Age	\
0	22	11628.0	3213	0.0	36	
1	22	30510.0	1117	0.0	75	
2	19	7450.0	3732	1.0	32	
3	19	17780.0	4378	0.0	52	
4	16	10404.0	2190	0.0	44	

	source
0	train
1	train
2	train
3	train
4	train

```
[ ]:
```

0.2 data cleaning

```
[18]: def na_randomfill(series):
    na_mask = pd.isnull(series)    # boolean mask for null values
    n_null = na_mask.sum()         # number of nulls in the Series

    if n_null == 0:
        return series              # if there are no nulls, no need to resample

    # Randomly sample the non-null values from our series
    # only sample this Series as many times as we have nulls
    fill_values = series[~na_mask].sample(n=n_null, replace=True,
    ↪random_state=0)

    # This ensures our new values will replace NaNs in the correct locations
    fill_values.index = series.index[na_mask]

    return series.fillna(fill_values)
```

filling null values with randomly same column values

```
[19]: data['Health Indicator'] = na_randomfill(data['Health Indicator'])
data['Holding_Policy_Duration'] = na_randomfill(data['Holding_Policy_Duration'])
data['Holding_Policy_Type'] = na_randomfill(data['Holding_Policy_Type'])
```

```
[20]: data.isnull().sum()
```


```
[20]: Accomodation_Type      0
      City_Code             0
      Health_Indicator       0
      Holding_Policy_Duration 0
      Holding_Policy_Type    0
      ID                    0
      Is_Spouse              0
      Lower_Age              0
      Reco_Insurance_Type    0
      Reco_Policy_Cat        0
      Reco_Policy_Premium    0
      Region_Code            0
      Response               21805
      Upper_Age              0
      source                 0
      dtype: int64
```

```
[21]: plt.subplot(131)
      data['Accomodation_Type'].value_counts().plot.
      ↳ bar(figsize=(24,6),title='Accomodation_Type')

      plt.subplot(132)
      data['Is_Spouse'].value_counts().plot.bar(title='Is_Spouse')

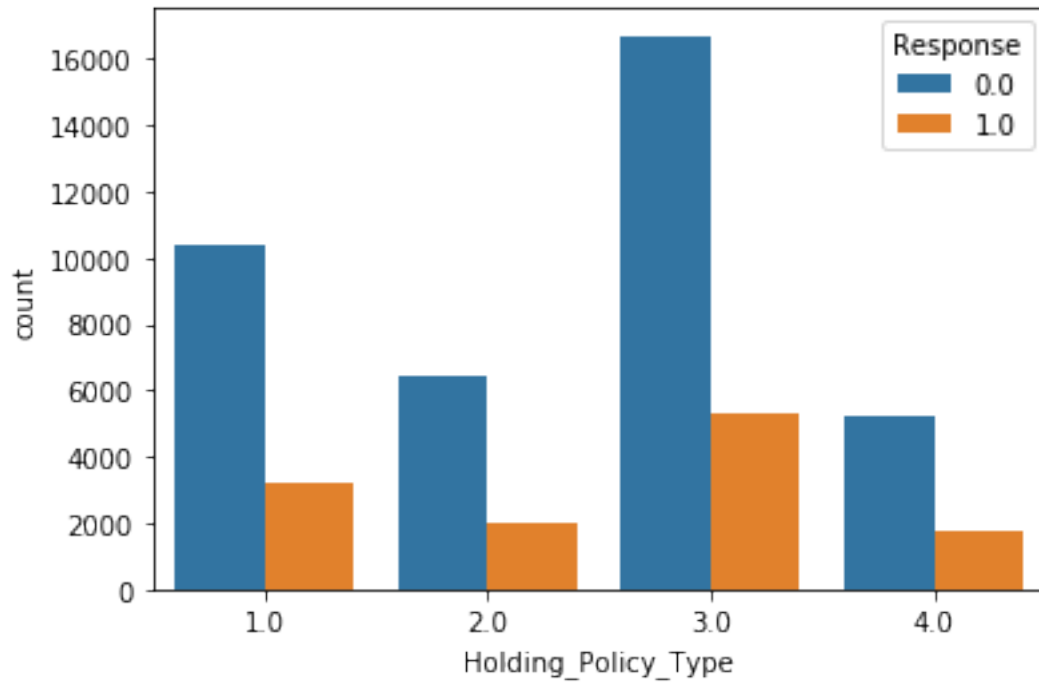
      plt.subplot(133)
      data['Reco_Insurance_Type'].value_counts().plot.bar(title='Reco_Insurance_Type')

      plt.show()
```



```
[22]: sns.countplot(hue='Response',x='Holding_Policy_Type',data=data)
```

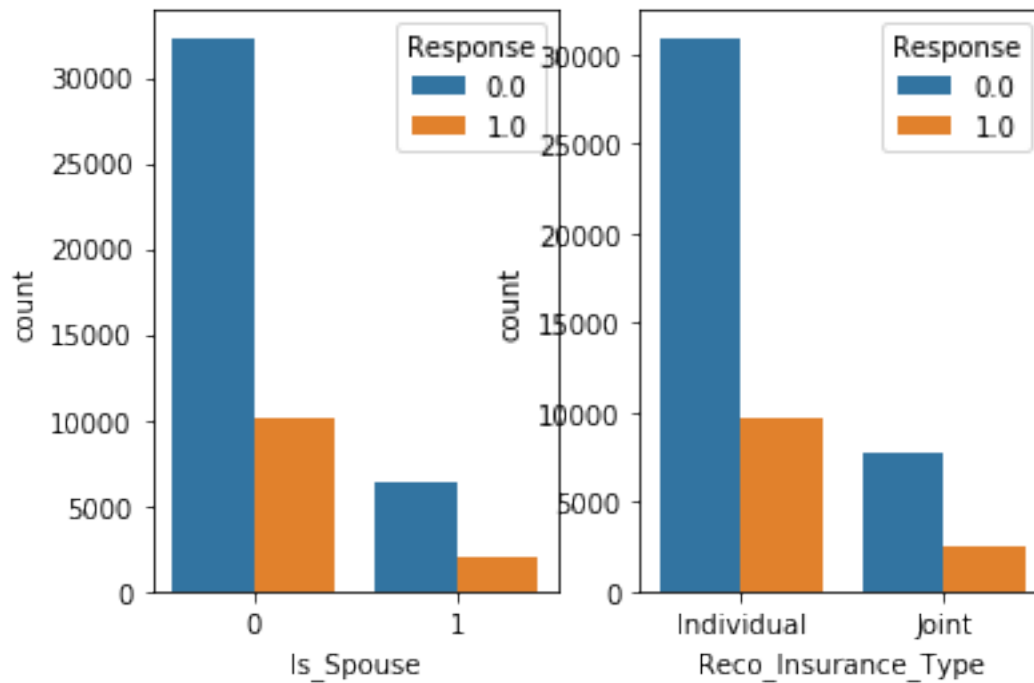
```
[22]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa23adc48>
```



```
[23]: plt.subplot(121)
sns.countplot(x='Is_Spouse',hue='Response',data=data)

plt.subplot(122)
sns.countplot(x='Reco_Insurance_Type',hue='Response',data=data)

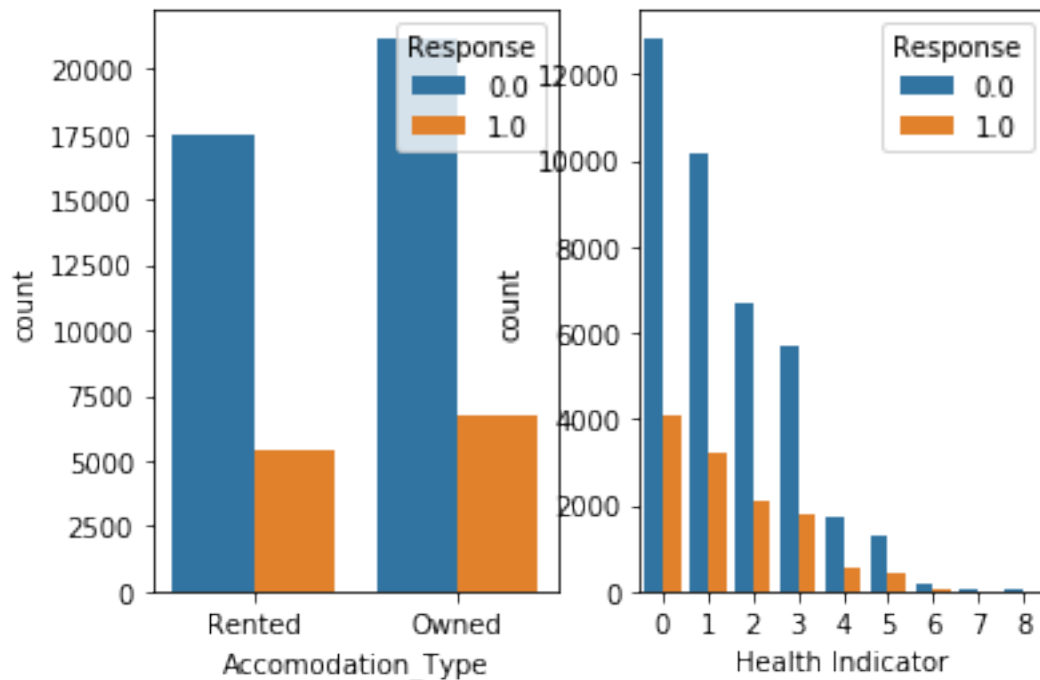
plt.show()
```

```
[24]: plt.subplot(121)
sns.countplot(x='Accommodation_Type', hue='Response', data=data)

plt.subplot(122)
sns.countplot(x='Health Indicator', hue='Response', data=data)

plt.show()
```

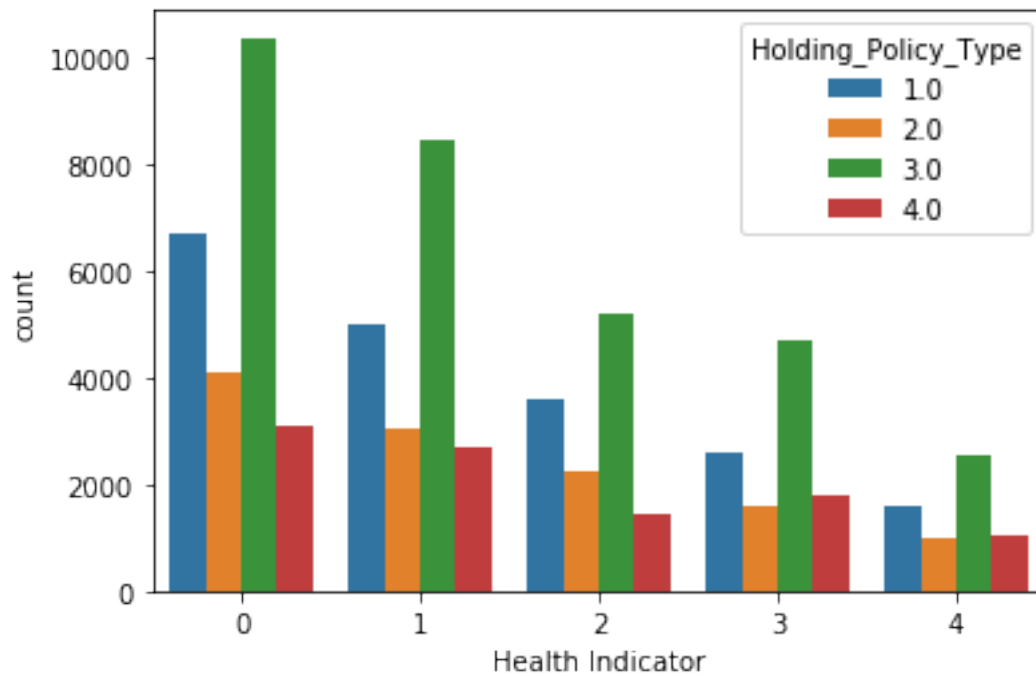


replacing values(5,6,7,8) with value(4).as you can see 5,6,7,8 didn't have much impact on response.

```
[25]: data['Health Indicator'].replace('7','4',inplace=True)
      data['Health Indicator'].replace('8','4',inplace=True)
      data['Health Indicator'].replace('5','4',inplace=True)
      data['Health Indicator'].replace('6','4',inplace=True)
```

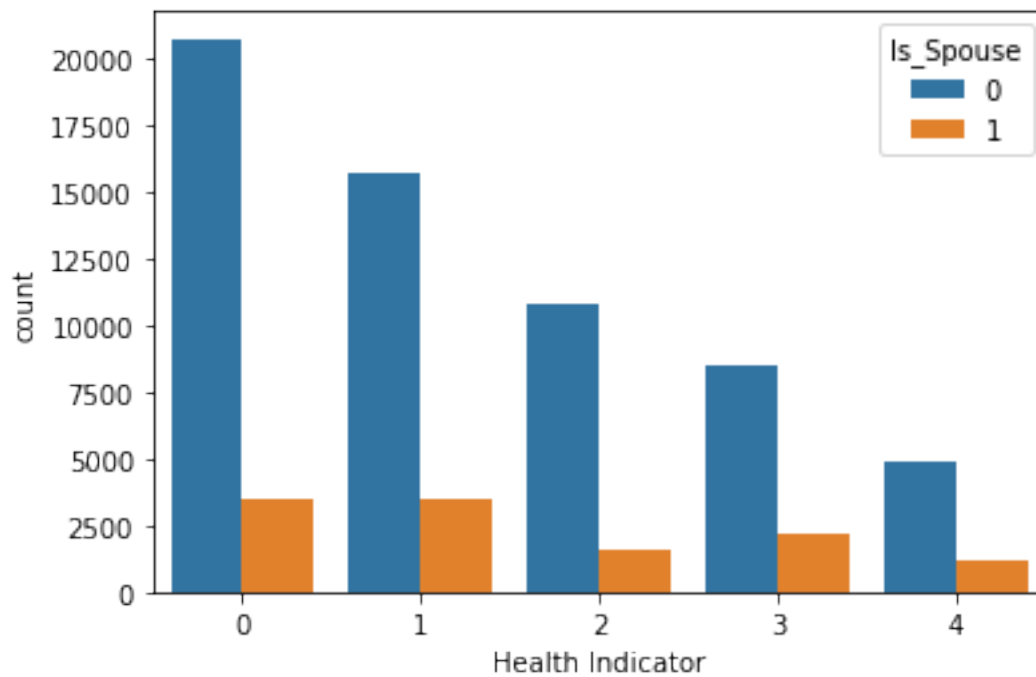
```
[26]: sns.countplot(x='Health Indicator',hue='Holding_Policy_Type',data=data)
```

```
[26]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa3916448>
```




```
[27]: sns.countplot(x='Health Indicator',hue='Is_Spouse',data=data)
```

```
[27]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa39be0c8>
```



```
[28]: g = sns.factorplot("City_Code", data=data, aspect=1.5, kind="count", color="b")
      g.set_xticklabels(rotation=75)
```

```
[28]: <seaborn.axisgrid.FacetGrid at 0x21aa39b65c8>
```



```
[29]: data.corr()
```

```
[29]:
```

	Holding_Policy_Type	ID	Is_Spouse	Lower_Age	\
Holding_Policy_Type	1.000000	0.006011	0.063155	0.068833	
ID	0.006011	1.000000	-0.003041	0.000758	
Is_Spouse	0.063155	-0.003041	1.000000	0.058470	
Lower_Age	0.068833	0.000758	0.058470	1.000000	
Reco_Policy_Cat	0.046751	0.000315	0.021489	0.020116	
Reco_Policy_Premium	0.074336	0.001245	0.510928	0.613374	
Region_Code	0.002804	0.004074	-0.002264	-0.004750	
Response	0.009312	0.005159	0.003859	-0.002099	
Upper_Age	0.079849	0.000066	0.198134	0.921175	


	Reco_Policy_Cat	Reco_Policy_Premium	Region_Code	\
Holding_Policy_Type	0.046751	0.074336	0.002804	
ID	0.000315	0.001245	0.004074	

Is_Spouse	0.021489	0.510928	-0.002264
Lower_Age	0.020116	0.613374	-0.004750
Reco_Policy_Cat	1.000000	0.060442	-0.062533
Reco_Policy_Premium	0.060442	1.000000	-0.013772
Region_Code	-0.062533	-0.013772	1.000000
Response	0.114321	0.007943	0.001121
Upper_Age	0.024325	0.791562	-0.006170

	Response	Upper_Age
Holding_Policy_Type	0.009312	0.079849
ID	0.005159	0.000066
Is_Spouse	0.003859	0.198134
Lower_Age	-0.002099	0.921175
Reco_Policy_Cat	0.114321	0.024325
Reco_Policy_Premium	0.007943	0.791562
Region_Code	0.001121	-0.006170
Response	1.000000	0.002772
Upper_Age	0.002772	1.000000

```
[30]: sns.heatmap(data.corr(), cmap="YlGnBu")
```


```
[30]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa3884dc8>
```



Let's check the distribution of Reco_Policy_Premium

```
[31]: sns.distplot(data['Reco_Policy_Premium'])
```

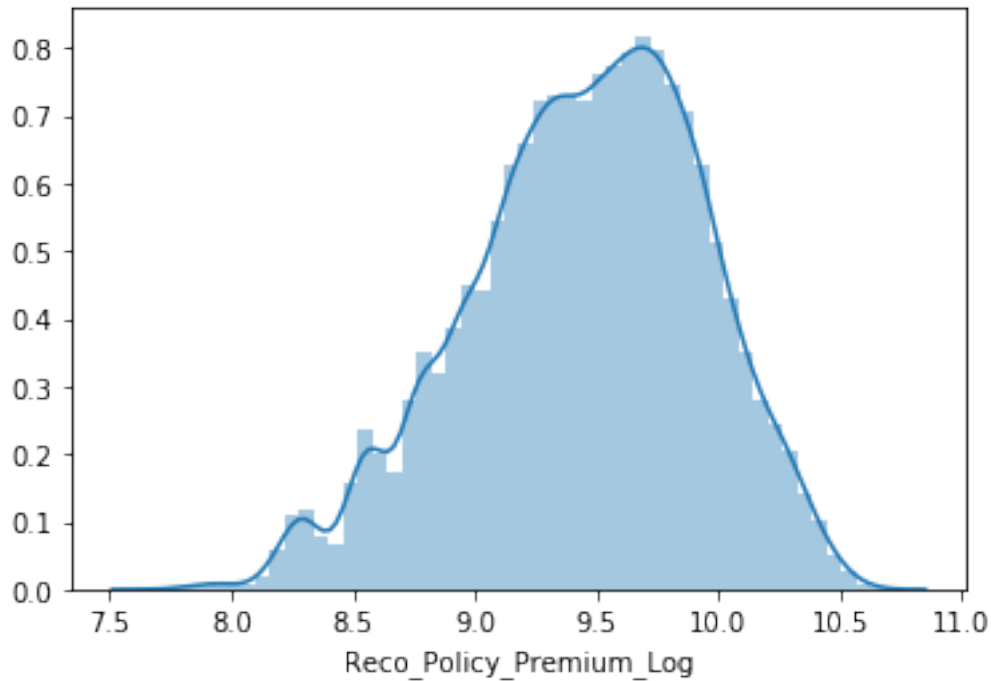
```
[31]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa38141c8>
```



we can see it shifted towards left, i.e.,the distribution is right skewed. so, let's take the log transformation to make the distribution normal.

```
[32]: data['Reco_Policy_Premium_Log'] = np.log(data['Reco_Policy_Premium'])  
sns.distplot(data['Reco_Policy_Premium_Log'])
```

```
[32]: <matplotlib.axes._subplots.AxesSubplot at 0x21aa24d7148>
```



Now the distribution looks much closer to normal and effect of extreme values has been significantly subsided.

```
[33]: data.dtypes
```

```
[33]: Accomodation_Type      object
      City_Code             object
      Health Indicator      object
      Holding_Policy_Duration  object
      Holding_Policy_Type     float64
      ID                    int64
      Is_Spouse              int64
      Lower_Age              int64
      Reco_Insurance_Type     object
      Reco_Policy_Cat         int64
      Reco_Policy_Premium     float64
      Region_Code            int64
      Response               float64
      Upper_Age              int64
      source                  object
      Reco_Policy_Premium_Log float64
      dtype: object
```

```
[34]: #One Hot InCoding:
```

```
data = pd.get_dummies(data, columns=['City_Code', 'Accommodation_Type', 'Health_Indicator', 'Holding_Policy_Type', 'Holding_Policy_Duration', 'Reco_Insurance_Type'])
```

```
[35]: data.shape
```

```
[35]: (72687, 73)
```

Let's convert data back into train and test data sets. Its generally a good idea to export both of these as modified data sets so that they can be re-used for multiple sessions. This can be achieved using following code:

```
[36]: train = data.loc[data['source']=="train"]
test = data.loc[data['source']=="test"]

#Drop unnecessary columns:
test.drop(['Response', 'source'], axis=1, inplace=True)
train.drop(['source', 'ID'], axis=1, inplace=True)

#Export files as modified versions:
train.to_csv("train_modified.csv", index=False)
test.to_csv("test_modified.csv", index=False)
```

0.3 Model Building

I would like to define a generic function which takes the algorithm and data as input and makes the model, performs roc_auc_score and generates submission.

```
[37]: #define target and ID column
target = 'Response'
IDcol = ['ID', 'Upper_Age']

from sklearn.metrics import roc_auc_score
from sklearn import metrics
def modelfit(alg, dtrain, dtest, pred, target, IDcol, filename):
    #fit the algorithm on the data
    alg.fit(dtrain[pred], dtrain[target])

    #predict training set
    dtrain_pred = alg.predict_proba(dtrain[pred])[:, 1]
    print("auc_score : %.4g" % roc_auc_score(dtrain[target].values, dtrain_pred))

    #Predict on testing data:
    dtest[target] = alg.predict_proba(dtest[pred])[:, 1]

    #Export submission file:
```



```
IDcol.append(target)
submission = pd.DataFrame({ x: dtest[x] for x in IDcol})
del submission['Upper_Age']
submission.to_csv(filename, index=False)
```

```
[41]: from lightgbm import LGBMClassifier
pred = [x for x in train.columns if x not in [target]+IDcol]
alg4 = LGBMClassifier(learning_rate = 0.01,
                      max_depth = 10,
                      num_leaves = 80,
                      n_estimators = 650)
modelfit(alg4, train, test, pred, target, IDcol, 'lgb1.csv')
```

auc_score : 0.7867

AV private Leaderboard score:0.6675

AV final rank:334

[]: